

# Semi-Supervised Classification with Graph Convolutional Networks

Thomas N. Kipf<sup>1</sup> & Max Welling<sup>1,2</sup>, <sup>1</sup>University of Amsterdam, <sup>2</sup>CIFAR

## End-to-end learning on graphs

Previous state-of-the-art for **classification on graphs**:

- **Graph kernels/embeddings** [1,2]: not trainable end to end
- **Regularization-based methods** [3-5]:

$$\mathcal{L} = \mathcal{L}_0 + \lambda \mathcal{L}_{\text{reg}}, \text{ with } \mathcal{L}_{\text{reg}} = \sum_{i,j} A_{ij} \|f(X_i) - f(X_j)\|^2$$

$A$ : adjacency matrix of graph  
 $X$ : matrix of node feature vectors

**Our approach:**

Learn graph-based model  $f(X, A)$  end to end.

## Fast graph convolutions

**Graph convolution** with signal  $x \in \mathbb{R}^N$  and filter  $g_\theta$ :

$$g_\theta \star x = U g_\theta(\Lambda) U^\top x$$

$U$ : eigenvectors of graph Laplacian  $L = I_N - D^{-\frac{1}{2}} A D^{-\frac{1}{2}}$

Comp. complexity:  $\mathcal{O}(N^2)$

**First-order approximation:**

$$g_\theta \star x \approx \theta_0 x + \theta_1 D^{-\frac{1}{2}} A D^{-\frac{1}{2}} x$$

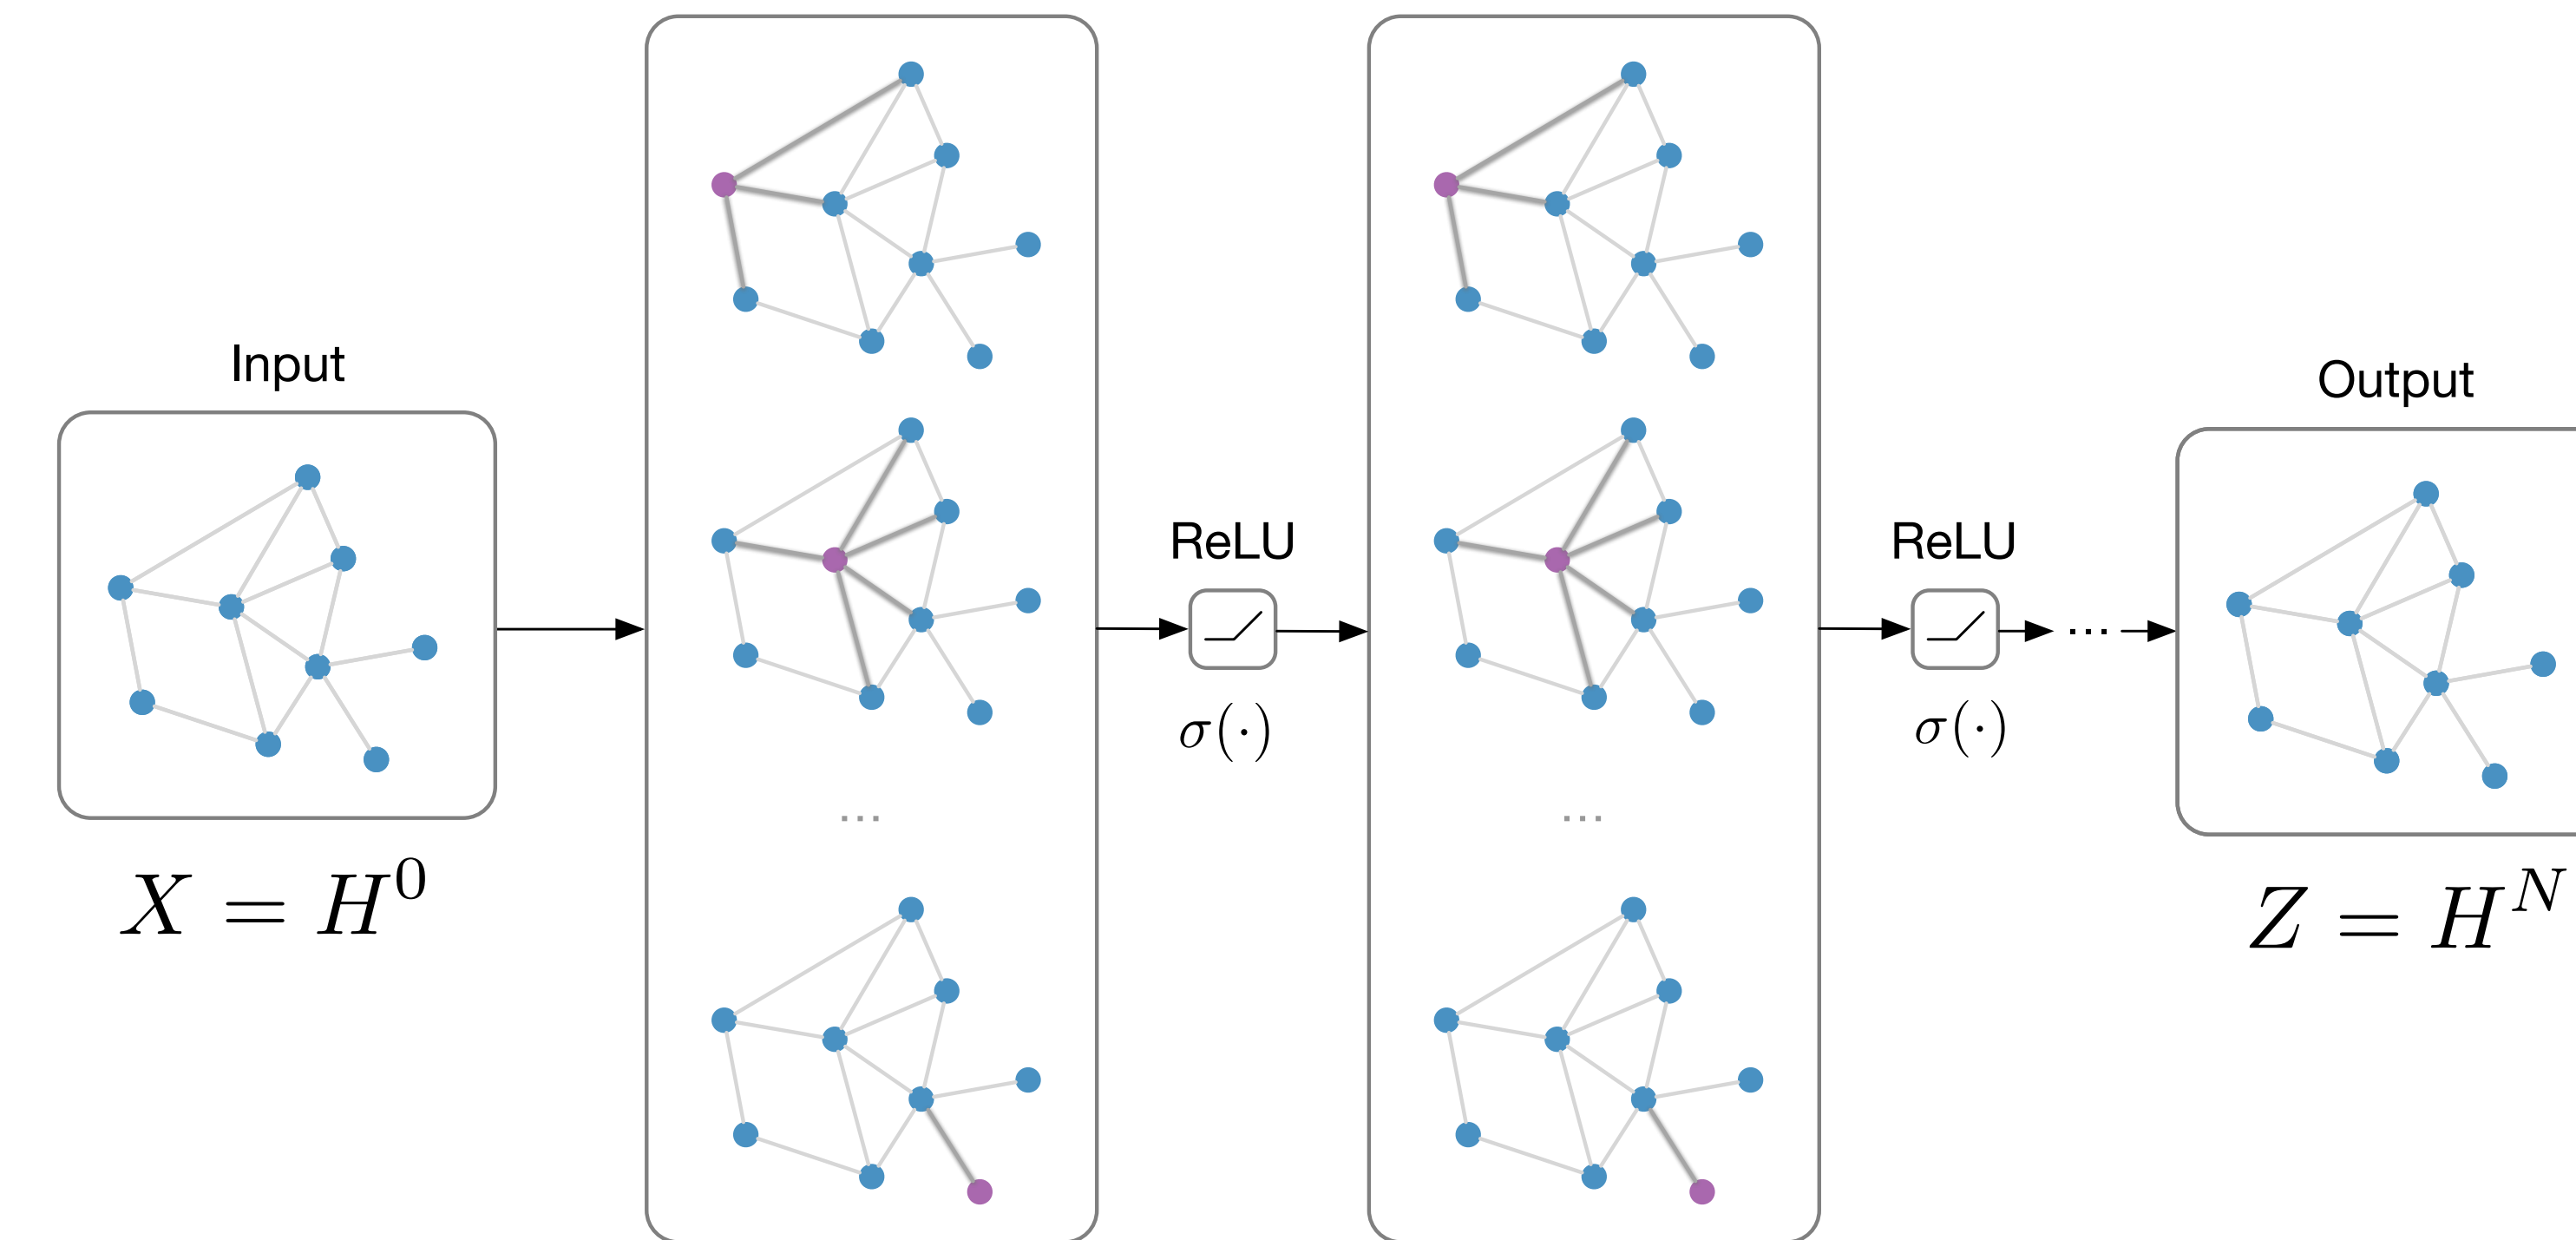
**Single parameter and renormalization\*:**

$$g_\theta \star x \approx \theta \bar{D}^{-\frac{1}{2}} \bar{A} \bar{D}^{-\frac{1}{2}} x$$

New complexity:  $\mathcal{O}(|\mathcal{E}|)$

$$* \bar{A} = A + I_N, \bar{D}_{ii} = \sum_j \bar{A}_{ij}$$

## Graph convolutional networks



$$H^{l+1} = \sigma(\hat{A} H^l W^l), \hat{A} = \bar{D}^{-\frac{1}{2}} \bar{A} \bar{D}^{-\frac{1}{2}}$$

$$\textbf{Two-layer model: } Z = f(X, A) = \text{softmax}(\hat{A} \text{ReLU}(\hat{A} X W^0) W^1)$$


## Experiments

**Goal:** Classify *all nodes* given only a few labels (test set: 1000 nodes)

Dataset	Type	Nodes	Edges	Classes	Features	Label rate
Citeseer	Citation network	3,327	4,732	6	3,703	0.036
Cora	Citation network	2,708	5,429	7	1,433	0.052
Pubmed	Citation network	19,717	44,338	3	500	0.003
NELL	Knowledge graph	65,755	266,144	210	5,414	0.001

Graph dataset statistics

## Graph embeddings



t-SNE embedding of hidden layer activations for Cora dataset (5% labels)

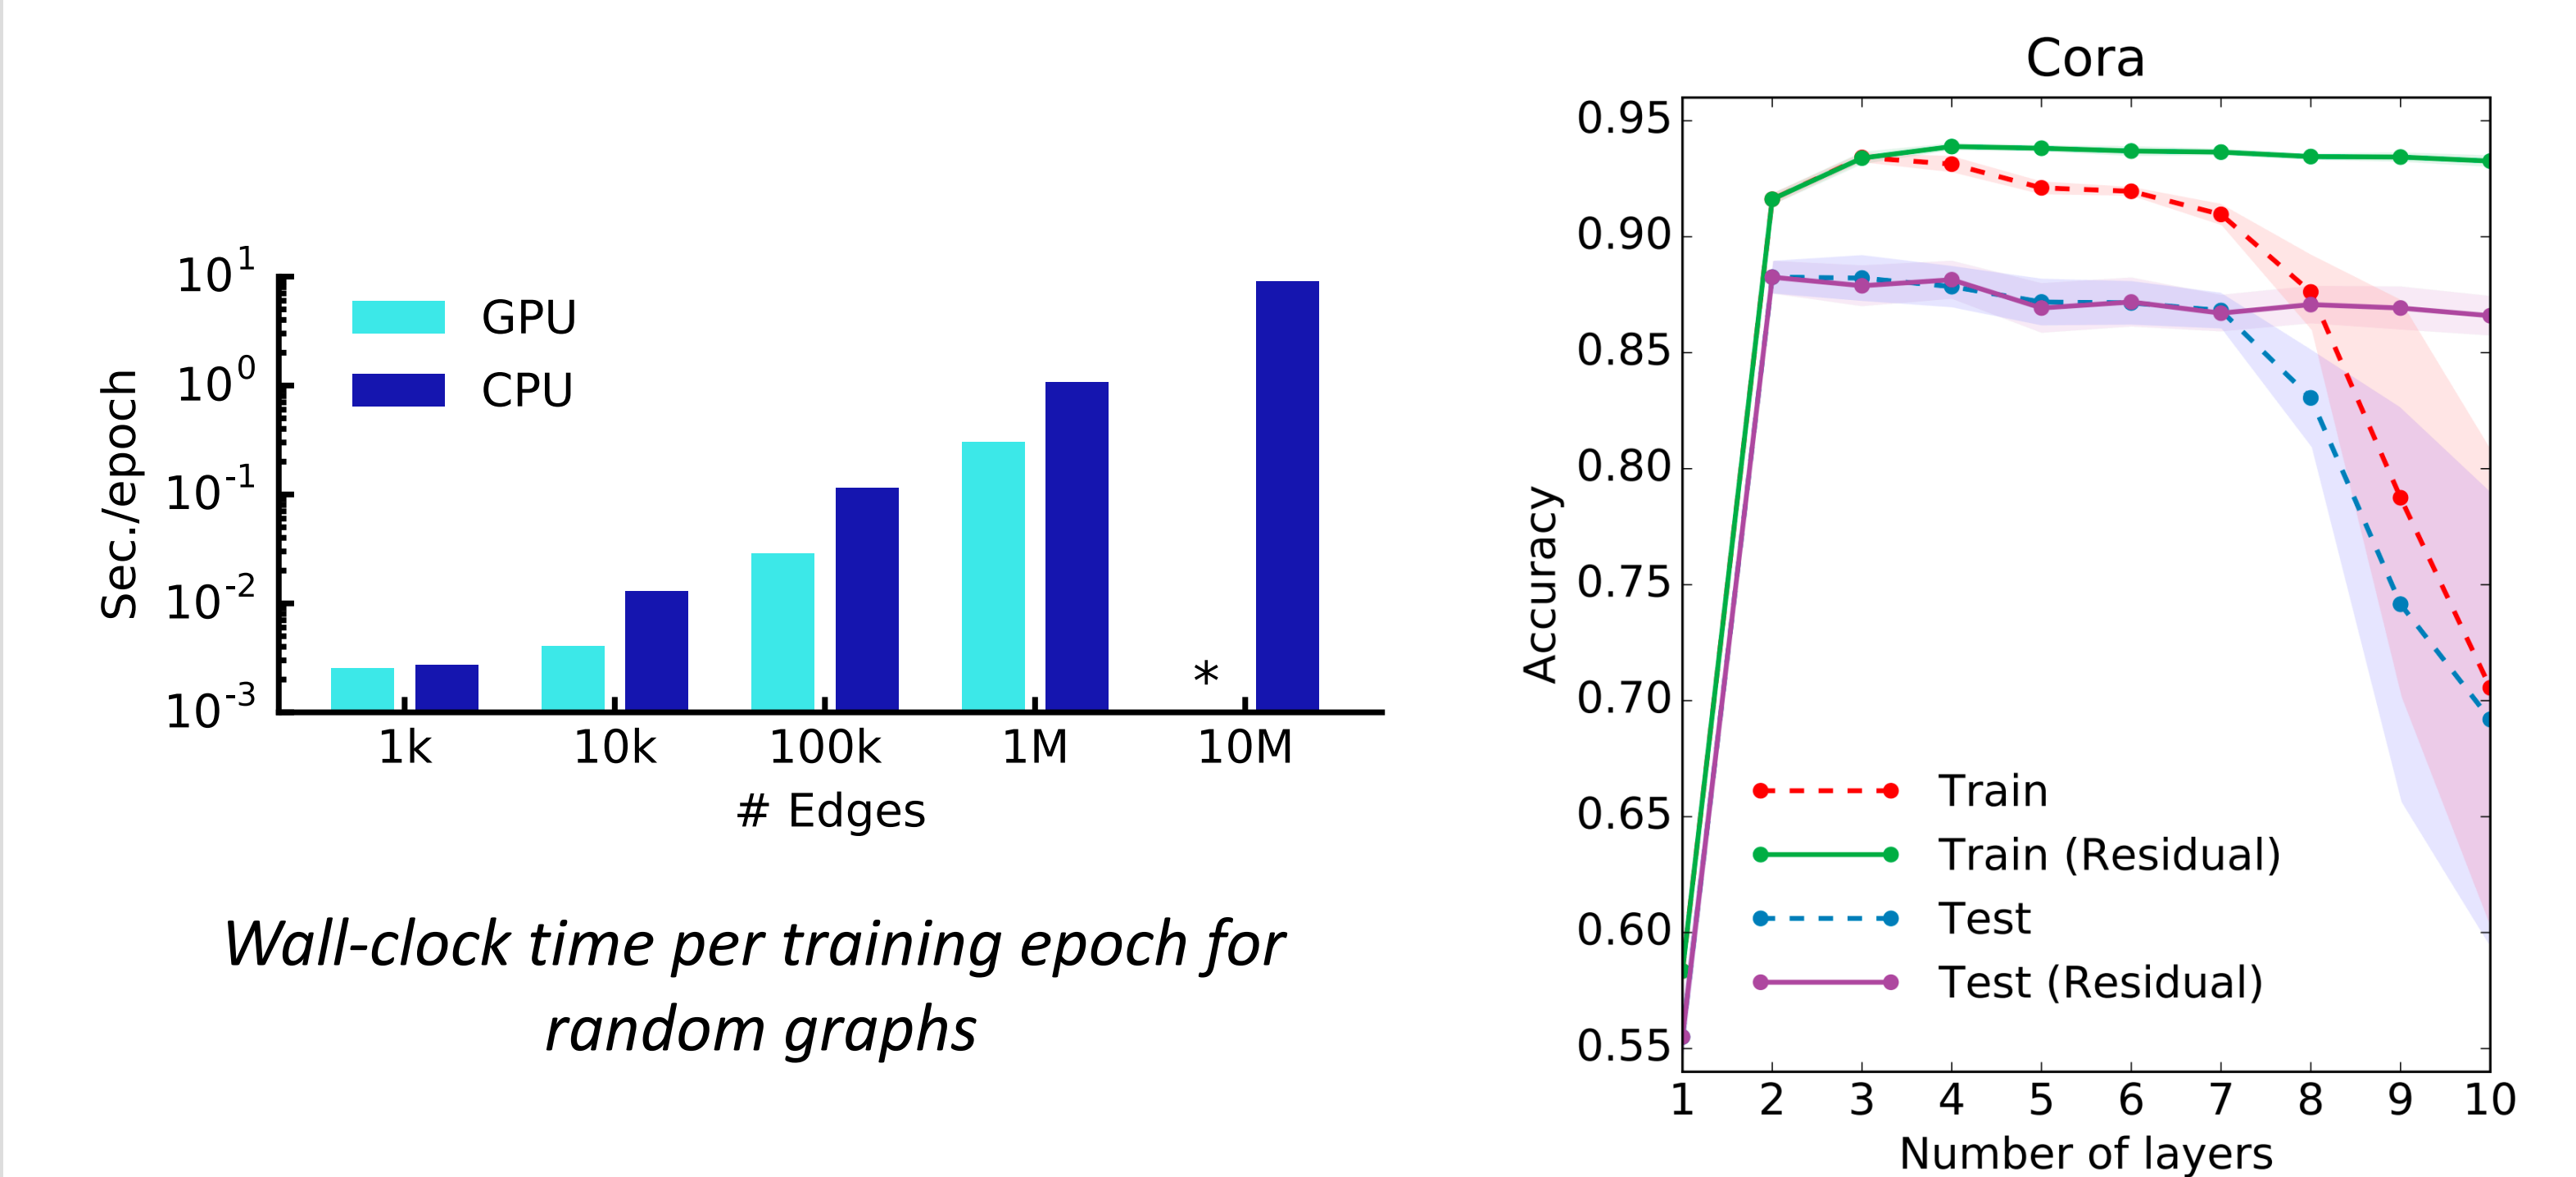
## Node classification results

Method	Citeseer	Cora	Pubmed	NELL
ManiReg [4]	60.1	59.5	70.7	21.8
SemiEmb [5]	59.6	59.0	71.1	26.7
LP [3]	45.3	68.0	63.0	26.5
DeepWalk [1]	43.2	67.2	65.3	58.1
Planetoid* [2]	64.7 (26s)	75.7 (13s)	77.2 (25s)	61.9 (185s)
<b>GCN (ours)</b>	<b>70.3 (7s)</b>	<b>81.5 (4s)</b>	<b>79.0 (38s)</b>	<b>66.0 (48s)</b>
GCN (rand. splits)	67.9 ± 0.5	80.1 ± 0.5	78.9 ± 0.7	58.4 ± 1.7

Model comparison (classification accuracy)

Description	Propagation model	Citeseer	Cora	Pubmed
Chebyshev filter [6]	$K=3$ $K=2$ $\sum_{k=0}^K T_k(\tilde{L}) X \Theta_k$	69.8 69.6	79.5 81.2	74.4 73.8
1 <sup>st</sup> -order model	$X \Theta_0 + D^{-\frac{1}{2}} A D^{-\frac{1}{2}} X \Theta_1$	68.3	80.0	77.5
Single parameter	$(I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}}) X \Theta$	69.3	79.2	77.4
<b>Renormalization trick</b>	$\bar{D}^{-\frac{1}{2}} \bar{A} \bar{D}^{-\frac{1}{2}} X \Theta$	<b>70.3</b>	<b>81.5</b>	<b>79.0</b>
1 <sup>st</sup> -order term only	$D^{-\frac{1}{2}} A D^{-\frac{1}{2}} X \Theta$	68.7	80.5	77.8
Multi-layer perceptron	$X \Theta$	46.5	55.1	71.4

Comparison of graph propagation models




Wall-clock time per training epoch for random graphs


Performance vs. model depth

## Message passing interpretation

Consider this graph:



Update node in red:



**Update rule:**

$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

$\mathcal{N}_i$ : neighbor indices  
 $c_{ij}$ : norm. constant (per edge)

## Conclusions

- End-to-end learning for node classification on graphs
- Scalable up to ~10M nodes/edges
- **Future work:** model variants (e.g. attention mechanism), edge features, scalability, unsupervised learning

**Code:** [github.com/tkipf/gcn](https://github.com/tkipf/gcn)

**Paper:** [arxiv.org/abs/1609.02907](https://arxiv.org/abs/1609.02907)

**Blog:** [tkipf.github.io/graph-convolutional-networks](https://tkipf.github.io/graph-convolutional-networks)



Project funded by SAP

### References:

- [1] Perozzi et al., Deepwalk: Online learning of social representations, SIGKDD (2014)  
[2] Yang et al., Revisiting semi-supervised learning with graph embeddings, ICML (2016)  
[3] Zhou et al., Learning with local and global consistency, NIPS (2004)

- [4] Belkin et al., Manifold regularization, JMLR (2006)  
[5] Weston et al., Deep learning via semi-supervised embedding, Neural Networks: Tricks of the Trade (2012)  
[6] Defferrard et al., Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering, NIPS (2016)

**Acknowledgements:** Christos Louizos, Taco Cohen, Rianne van den Berg, Karen Ullrich, Joan Bruna, Zhilin Yang, Dave Herman, Pramod Sinha, and Abdul-Saboor Sheikh

**Web:** [tkipf.github.io](https://tkipf.github.io) — [staff.fnwi.uva.nl/m.welling](mailto:staff.fnwi.uva.nl/m.welling)  
**Mail:** {t.n.kipf, m.welling}@uva.nl